MeetingDNA

Decision Intelligence Engine

"Git for decisions — every choice your company made, who made it, why, and what broke because of it."

Hackathon Project Documentation

B2B SaaS — Decision Intelligence

2025

Table of Contents

Table of Contents 2
1. Executive Summary 3
2. Problem Statement 3
  2.1 The Institutional Amnesia Problem 3
  2.2 Why Existing Tools Fail 3
  2.3 The Real Cost 4
3. Proposed Solution 4
  3.1 What MeetingDNA Builds 4
  3.2 The Decision Node Structure 4
  3.3 Core Capabilities 5
  3.4 The Demo Moment 5
4. Technical Architecture 6
  4.1 Tech Stack Overview 6
  4.2 Pipeline Architecture 6
  4.3 Why This Stack Is Novel 7
5. Market Analysis 7
  5.1 Market Opportunity 7
  5.2 Competitive Positioning 8
6. Business Model 8
  6.1 Pricing Tiers 8
  6.2 Revenue Moats 9
7. Idea Analysis & Scorecard 9
  7.1 Evaluation Scorecard 9
  7.2 Key Strengths 9
  7.3 Risks & Mitigations 10
8. Official Submission Abstract 10
9. Appendix — Sample Query Demonstrations 11
  Query 1: Decision Retrieval 11
  Query 2: Dependency Graph Traversal 11
  Query 3: Owner Accountability 12

1. Executive Summary

MeetingDNA is a B2B SaaS decision intelligence engine that automatically converts company meeting recordings and transcripts into a structured, queryable decision knowledge graph — with zero behavior change required from any team member.

It answers the question every growing company struggles with: 'Why did we decide that, and what else broke when we changed it?'

Every company above ten people silently suffers from institutional amnesia. Decisions get made in meetings, reasoning stays locked in someone's head, and months later nobody can explain why the pricing changed, who owned the infra migration, or what assumptions the product roadmap was built on.

Existing tools like Notion and Confluence attempt to solve this through documentation — but documentation requires discipline that teams consistently fail to maintain. MeetingDNA takes a fundamentally different approach: it works on the recordings your team already makes, requiring nothing new from anyone.

2. Problem Statement

2.1 The Institutional Amnesia Problem

Most organizations operate with a painful but invisible inefficiency: they have no reliable record of why decisions were made. The consequences compound quietly across every team, every quarter.

- Decisions get reversed without anyone knowing what downstream dependencies they break
- New hires spend weeks reconstructing context that was never written down
- Engineering teams build on assumptions that were invalidated in meetings they never attended
- Strategy pivots create cascading misalignments nobody tracks or surfaces

2.2 Why Existing Tools Fail

Tool	What It Does & Why It Falls Short

Notion / Confluence	Requires manual documentation. Nobody writes down the 'why.' Adoption fails within weeks.
Gong / Chorus	Sales call intelligence only. No decision graph. No org-level memory.
Fireflies.ai	Transcription and keyword search. Returns text chunks, not structured decisions.
Slack / Email	Conversational. Context buried and unsearchable after two weeks.
Mem.ai / Rewind	Personal knowledge tools. Individual scope, not organizational.

The gap is structural: every tool in the market captures what was said. Nobody captures what was decided, who owns it, what rationale drove it, and what other decisions depend on it. That gap is exactly where MeetingDNA operates.

2.3 The Real Cost

This is not a minor inconvenience. Consider a Series B company of 80 people running two hours of meetings per day. Over six months that's roughly 15,000 hours of decisions made — none of them indexed, none of them queryable, none of them connected to each other. When a single upstream decision reverses and three downstream decisions were built on it, the misaligned engineering work alone can represent hundreds of thousands of dollars in waste.

3. Proposed Solution

3.1 What MeetingDNA Builds

MeetingDNA is not a meeting summarizer. It is not RAG over transcripts. It is a structured decision knowledge graph — where every node represents a decision with five properties: what was decided, who owns it, why it was made, its current status, and which other decisions depend on it.

3.2 The Decision Node Structure

Property	Description	Example Value	Why It Matters
Decision	The actual choice made	Move infra to AWS by Q4	Core record of intent

Owner	Accountable individual or team	Rohan, Engineering Lead	Accountability layer
Rationale	Why this choice was made	GCP costs 40% higher at scale	Context preservation
Status	Current state of the decision	Reversed — Nov 2024 meeting	Temporal accuracy
Dependencies	Downstream decisions built on this	3 linked decisions affected	Impact intelligence

3.3 Core Capabilities

Natural Language Query Interface

- 'Why did we drop the freemium tier?' → sourced answer with original meeting timestamp
- 'Who owned the AWS migration?' → owner, decision date, current status
- 'What changed in Q3 pricing?' → full decision history with rationale chain

Dependency Graph Traversal

- 'Show all decisions that depended on the infra migration' → graph lights up with linked nodes
- When a decision reverses, MeetingDNA surfaces every downstream decision never revisited
- Full causal chain visualization — understand second-order impacts instantly

Retroactive Intelligence

- Works on recordings and transcripts your team already has — no new behavior required
- Ingests Zoom recordings, Gong transcripts, Fireflies exports, Otter.ai outputs
- Builds institutional memory going back to day one of your recording history

3.4 The Demo Moment

The single most powerful demo sequence:

1. Drop 5 company transcripts covering 6 months of decisions
2. MeetingDNA builds the decision graph — live, in front of the judges
3. Query: 'Why did we drop freemium?' → sourced answer with meeting clip reference

4. Query: 'What decisions depended on the AWS migration?' → graph illuminates 3 nodes
5. **Punch line: 'Two of those decisions were never revisited after the migration reversed. That is $200K of misaligned engineering work.'**

4. Technical Architecture

4.1 Tech Stack Overview

Layer	Technology & Role
Transcription	OpenAI Whisper — converts audio recordings to raw text transcripts
Decision Extraction	Claude API (Anthropic) — long-context LLM that extracts structured decision nodes: owner, rationale, status, dependencies as JSON
Graph Storage	Neo4j — stores decision nodes and their relationships; enables graph traversal queries
Semantic Search	ChromaDB — vector embeddings for semantic fallback when graph traversal is insufficient
Backend Orchestration	FastAPI — API layer that coordinates the extraction pipeline and query agent
Demo Interface	Streamlit — clean, functional UI with graph visualization for the demo

4.2 Pipeline Architecture

```
INPUT LAYER
  Audio Files / Existing Transcripts / Gong Exports / Fireflies Outputs
          |
TRANSCRIPTION LAYER
  OpenAI Whisper  →  Raw Text Transcripts
          |
EXTRACTION LAYER
  Claude API (Long Context)  →  Structured Decision JSON
  { decision, owner, rationale, status, dependencies[] }
          |
STORAGE LAYER
```

```
  Neo4j Graph DB  ⟷  ChromaDB Vector Store
  (Relational Traversal)    (Semantic Search Fallback)
         |
QUERY AGENT
  FastAPI  →  Claude API  →  Natural Language Response
         |
INTERFACE LAYER
  Streamlit UI  +  Neo4j Graph Visualization
```

4.3 Why This Stack Is Novel

Most hackathon projects follow a predictable pattern:

```
Data  →  Embedding  →  Vector DB  →  RAG  →  Chat UI
```

MeetingDNA's pipeline is architecturally distinct:

```
Transcripts  →  LLM Extraction  →  Structured Decision Nodes
            →  Graph DB  →  Relational + Semantic Query  →  Agent Response
```

The three layers that make this genuinely novel are:

- **Relation Extraction:** Relation extraction — pulling owner, rationale, status, and dependency links from unstructured conversation is a non-trivial NLP task that most teams will not attempt
- **Graph Traversal:** Graph traversal over decisions — querying 'what broke when this decision reversed' requires graph logic, not vector similarity. This is architecturally distinct from RAG
- **Temporal Reasoning:** Temporal reasoning — decisions have status that evolves over time. Tracking that state change across six months of meetings is genuinely hard

---

5. Market Analysis

5.1 Market Opportunity

Market Level	Definition	Size Estimate	Notes
TAM	Every B2B company with recorded meetings	Global knowledge economy	Essentially every SaaS company
SAM	Mid-market SaaS, 50-500 employees, already using Zoom/Gong/Notion	~500,000 companies globally	Primary target segment
SOM	10,000 companies × $500/month	$60M ARR	Conservative 5-year penetration

5.2 Competitive Positioning

MeetingDNA occupies the whitespace between Gong and Notion — the decision layer that sits between sales intelligence and documentation, and is owned by nobody.

Competitor	Their Category	Their Gap	MeetingDNA Advantage
Gong	Sales call intelligence	Sales only, no decision graph	Org-wide, all departments
Fireflies.ai	Meeting transcription	Text search only, no structure	Structured graph, not search
Notion AI	Docs and knowledge base	Requires manual documentation	Automatic, retroactive
Mem.ai	Personal knowledge graph	Individual scope only	Organizational scope
Rewind	Personal meeting memory	Personal, not team-level	Team and org intelligence

6. Business Model

6.1 Pricing Tiers

Tier	Price	Target	Key Features
Starter	$49/seat/month	Startups, 10-50 employees	Up to 50 hours of transcripts, basic graph, NL queries
Growth	$199/month flat	Series A-B, 50-200 employees	Unlimited transcripts, full graph, integrations, team dashboard

Enterprise	Custom pricing	500+ employees	On-premise option, SSO, custom integrations, SLA

6.2 Revenue Moats

- **Integration Lock-in:** Integration depth — Zoom, Gong, Fireflies, Slack, Notion integrations compound switching cost with every connection added
- **Data Network Effect:** Graph depth — by 12 months, a company's entire decision history lives in MeetingDNA. The switching cost becomes existential
- **Passive Value Creation:** Zero adoption friction — because no behavior change is required, churn from 'we stopped using it' is structurally near-impossible

7. Idea Analysis & Scorecard

7.1 Evaluation Scorecard

Dimension	Score	Rating
Problem Validity	9/10	★★★★★
Market Opportunity	8/10	★★★★☆
Technical Novelty	9/10	★★★★★
Demo-ability	10/10	★★★★★
Business Model Clarity	8/10	★★★★☆
Competitive Position	7/10	★★★★☆
Execution Risk (Hackathon)	6/10	★★★☆☆
Team Fit	9/10	★★★★★
Overall Score	8.25/10	★★★★★

7.2 Key Strengths

- The problem is universally felt — every person in a B2B company has experienced 'why did we decide this?' pain
- No direct competitor — this is a category creation play, not a feature war

- Demo has emotional resonance — the '$200K of misaligned work' punch line lands with every judge, investor, and founder in the room
- ML background creates authentic technical depth — entity extraction and relation classification are not faked skills here
- Zero behavior change adoption — structurally eliminates the most common reason enterprise software fails

7.3 Risks & Mitigations

Risk	Mitigation Strategy
Pipeline complexity — multiple failure points under time pressure	Pre-transcribe demo audio. Pre-populate part of the graph. Have a static fallback demo ready.
'Mistaken for meeting summarizer' positioning	Lead every pitch with the dependency graph demo, not the Q&A feature. The graph visual is your differentiator.
LLM extraction quality on noisy transcripts	Use curated demo transcripts. Show it works cleanly and don't let judges probe edge cases live.
Neo4j setup time in hackathon	Use Neo4j Aura (cloud) not local. Or fall back to NetworkX for graph logic if time is tight.

8. Official Submission Abstract

The following is the polished abstract for hackathon idea submission:

PROBLEM STATEMENT

Every growing company suffers from institutional amnesia. Decisions get made in meetings, rationale lives in someone's head, and six months later nobody remembers why the pricing changed or who owned the infra migration. Tools like Notion and Confluence try to fix this through documentation — but documentation requires discipline that teams never sustain. The result is reversed decisions, repeated mistakes, and weeks lost reconstructing context that was never written down.

PROPOSED SOLUTION

MeetingDNA converts your existing meeting recordings into a structured decision intelligence layer — automatically, retroactively, with zero behavior change required. It extracts every decision made across meetings along with the owner, rationale, current status, and downstream dependencies, storing them as a traversable decision graph. Anyone can then query it in plain English. The key differentiator is dependency tracking — when a decision reverses, MeetingDNA immediately surfaces every downstream decision built on top of it and never revisited.

TECH STACK

Whisper (transcription) · Claude API (extraction + agent) · Neo4j (decision graph) · ChromaDB (semantic search) · FastAPI (backend) · Streamlit (demo interface)

9. Appendix — Sample Query Demonstrations

Query 1: Decision Retrieval

```
User: "Why did we drop the freemium tier?"


MeetingDNA: The freemium tier was discontinued on September 14, 2024. The decision
was owned by Priya (Head of Growth). Rationale: conversion rate from free to paid
was 0.8% vs industry benchmark of 3-5%. The engineering cost of maintaining the
free tier was consuming 20% of backend capacity. Source: Product Strategy Sync,
Sep 14 2024, timestamp 00:23:41.
```

Query 2: Dependency Graph Traversal

```
User: "What decisions depended on the AWS migration?"


MeetingDNA: The AWS migration decision (Aug 2024, Owner: Rohan) has 3 downstream
dependencies:
  1. CDN vendor switch to CloudFront [Completed — aligned]
  2. Database replication strategy overhaul [In progress — aligned]
```

```
  3. Pricing model update for compute costs [NEVER REVISITED — decision made
     before migration reversed. This decision is now based on outdated assumptions.]

Warning: Decision #3 was built on assumptions that may no longer hold.
Recommend immediate review.
```

Query 3: Owner Accountability

```
User: "Who owns the infra migration and what is its current status?"

MeetingDNA: The infra migration to AWS was owned by Rohan (Engineering Lead).
Status: REVERSED on November 8, 2024.
Original rationale: Cost reduction and better ML tooling ecosystem.
Reversal reason: Integration complexity exceeded estimated timeline by 3x.
Source: Engineering All-Hands, Nov 8 2024, timestamp 01:12:05.
```

MeetingDNA — Built for the builders who refuse to repeat history.

© 2025 MeetingDNA. All rights reserved.